[SOLVED] Intel NIC e1000e hardware unit hang

md127 · Mar 5, 2022

Forums > Proxmox Virtual Environment > **Proxmox VE: Installation and configuration**

M
md127
New Member

Mar 5, 2022 #1

This has been discussed many times, is there a solution/workaround yet, as the specified steps are not working for me?
I've the following in /etc/network/interfaces

Code:

```
auto lo
iface lo inet loopback

iface eno1 inet manual
        offload-gso off
        offload-gro off
        offload-tso off
        offload-rx off
        offload-tx off
        offload-rxvlan off
        offload-txvlan off
        offload-sg off
        offload-ufo off
        offload-lro off
auto vmbr0
iface vmbr0 inet static
        address xxx
        gateway xxx
        bridge-ports eno1
        bridge-stp off
        bridge-fd 0
```

Syslog:

Code:

```
Mar 05 09:38:22 snuc2 kernel: e1000e 0000:00:1f.6 eno1: Detected Hardware Ur
  TDH                  <3b>
  TDT                  <59>
  next_to_use          <59>
  next_to_clean        <3a>
buffer_info[next_to_clean]:
  time_stamp           <1000bb386>
  next_to_watch        <3b>
```

```
  jiffies               <1000bb5e0>
  next_to_watch.status <0>
MAC Status              <40080083>
PHY Status              <796d>
PHY 1000BASE-T Status   <3800>
PHY Extended Status     <3000>
PCI Status              <10>
Mar 05 09:38:24 snuc2 kernel: e1000e 0000:00:1f.6 eno1: Detected Hardware Ur
  TDH                   <3b>
  TDT                   <59>
  next_to_use           <59>
  next_to_clean         <3a>
buffer_info[next_to_clean]:
  time_stamp            <1000bb386>
  next_to_watch         <3b>
  jiffies               <1000bb7d8>
  next_to_watch.status <0>
MAC Status              <40080083>
PHY Status              <796d>
PHY 1000BASE-T Status   <3800>
PHY Extended Status     <3000>
PCI Status              <10>
Mar 05 09:38:26 snuc2 kernel: e1000e 0000:00:1f.6 eno1: Detected Hardware Ur
  TDH                   <3b>
  TDT                   <59>
  next_to_use           <59>
  next_to_clean         <3a>
buffer_info[next_to_clean]:
  time_stamp            <1000bb386>
  next_to_watch         <3b>
  jiffies               <1000bb9c8>
```

---

oguz
Proxmox Retired Staff
Retired Staff

Mar 14, 2022 #2

hi,

> md127 said:
>
> This has been discussed many times, is there a solution/workaround yet, as the specified steps are not working for me?

where did you find those steps?

could you install `ethtool` and check your interface like: `ethtool -k eno1 | grep offload` to see if they're really disabled?

if not then please check the workaround in the post here [0] (the original linked thread is [1] but quite a long one to go through), using ethtool to set those options off.

if none of that works, post the requested information from my first post on thread [0]

[0]: https://forum.proxmox.com/threads/trap-error-on-e1000-network-adapter.105758
[1]: https://forum.proxmox.com/threads/e1000-driver-hang.58284/page-4#post-302307

Best regards,
Oguz

Do you already have a Commercial Support Subscription? - If not, Buy now and read the documentation

---

M

md127
New Member

Mar 14, 2022 #3

> oguz said:
>
> could you install `ethtool` and check your interface like: `ethtool -k eno1 | grep offload` to see if they're really disabled?

@oguz That's a great tip as although I had them disabled in the interface file (see below), the offloading got enabled automatically! Do we need to setup a cron job to periodically disable this?

> Code:
>
> ```
> auto lo
> iface lo inet loopback
>
> iface eno1 inet manual
>  offload-gso off
>  offload-gro off
>  offload-tso off
>  offload-rx off
>  offload-tx off
>  offload-rxvlan off
>  offload-txvlan off
>  offload-sg off
>  offload-ufo off
>  offload-lro off
> auto vmbr0
> iface vmbr0 inet static
>  address 192.168.1.158/24
>  gateway 192.168.1.1
>  bridge-ports eno1
>  bridge-stp off
>  bridge-fd 0
> ```

Mar 14, 2022 #4

oguz
Proxmox Retired Staff
Retired Staff

> md127 said:
>
> ! Do we need to setup a cron job to periodically disable this?

check the linked post [0], you can add a `post-up` directive in your network configuration file as described there (that will basically rerun the workaround when the interface is up)

Best regards,
Oguz

Do you already have a Commercial Support Subscription? - If not, Buy now and read the documentation

---

M
md127
New Member

Mar 14, 2022 #5

> oguz said:
>
> check the linked post [0], you can add a `post-up` directive in your network configuration file as described there (that will basically rerun the workaround when the interface is up)

Thanks! Apply them to the vmbr0 interface as well?
And reboot PVE?

---

oguz
Proxmox Retired Staff
Retired Staff

Mar 14, 2022 #6

> md127 said:
>
> Apply them to the vmbr0 interface as well?
> And reboot PVE?

yes and yes. you can also install `ifupdown2` and do `ifreload -a` instead of rebooting

Best regards,
Oguz

Do you already have a Commercial Support Subscription? - If not, Buy now and read the documentation

---

M
md127
New Member

Mar 14, 2022 #7

Awesome! thanks again. Will try and mark this as a "solution" if it works
Last question (for the time being): only gso/tso off? What about the others?
Right now I've switched off a lot more options in the interfaces file shared above, which gives me this:

> Code:

```
root@snuc2:~# ethtool -k eno1|grep offload
tcp-segmentation-offload: off
generic-segmentation-offload: off
generic-receive-offload: off
large-receive-offload: off [fixed]
rx-vlan-offload: off
tx-vlan-offload: off
l2-fwd-offload: off [fixed]
hw-tc-offload: off [fixed]
esp-hw-offload: off [fixed]
esp-tx-csum-hw-offload: off [fixed]
rx-udp_tunnel-port-offload: off [fixed]
tls-hw-tx-offload: off [fixed]
tls-hw-rx-offload: off [fixed]
macsec-hw-offload: off [fixed]
hsr-tag-ins-offload: off [fixed]
hsr-tag-rm-offload: off [fixed]
hsr-fwd-offload: off [fixed]
hsr-dup-offload: off [fixed]
```

---

oguz
Proxmox Retired Staff
Retired Staff

Mar 15, 2022 #8

md127 said:

Last question (for the time being): only gso/tso off? What about the others?

try the options as posted before: `ethtool -K eno1 gso off gro off tso off tx off rx off`

if it doesn't solve your issue please report

Best regards,
Oguz

Do you already have a Commercial Support Subscription? - If not, Buy now and read the documentation

---

M
md127
New Member

Mar 19, 2022 #9

@oguz guess that's it! I don't see any more hardware hangs. Thanks much for helping debug the issue and provide the workaround.
Is it possible to disable rx-vlan/tx-vlan offload in that script as well?

Code:

```
root@snuc2:~# ethtool -k eno1|grep offload
tcp-segmentation-offload: off
generic-segmentation-offload: off
generic-receive-offload: off
```

```
large-receive-offload: off [fixed]
rx-vlan-offload: off
tx-vlan-offload: off
l2-fwd-offload: off [fixed]
hw-tc-offload: off [fixed]
esp-hw-offload: off [fixed]
esp-tx-csum-hw-offload: off [fixed]
rx-udp_tunnel-port-offload: off [fixed]
tls-hw-tx-offload: off [fixed]
tls-hw-rx-offload: off [fixed]
macsec-hw-offload: off [fixed]
hsr-tag-ins-offload: off [fixed]
hsr-tag-rm-offload: off [fixed]
hsr-fwd-offload: off [fixed]
hsr-dup-offload: off [fixed]
```

haarm123

You must log in or register to reply here.

Share:

Forums > Proxmox Virtual Environment > **Proxmox VE: Installation and configuration**

About

The Proxmox community has been around for many years and offers help and support for Proxmox VE, Proxmox Backup Server, and Proxmox Mail Gateway.
We think our community is one of the best thanks to people like you!

Quick Navigation

Home

Get Subscription

Wiki

Downloads

Proxmox Customer Portal

About

Get your subscription!

The Proxmox team works very hard to make sure you are running the best software and getting stable updates and security enhancements, as well as quick enterprise support. Tens of

thousands of happy customers have a Proxmox subscription. Get your own in 60 seconds.

Buy now!

Community platform by XenForo® © 2010-2023 XenForo Ltd.

This site uses cookies to help personalise content, tailor your experience and to keep you logged in if you register.
By continuing to use this site, you are consenting to our use of cookies.

Accept

Learn more...